Chenru Wen

chenru.wen@mbzuai.ac.ae | wcr21.github.io

Research Interests

I am interested in how robots should act when key information is not directly observed: an occluded target, a slipped object, or a hidden mass. My work is at the intersection of robot learning, manipulation, and state estimation.

Education

Mohamed bin Zayed University of Artificial Intelligence (MBZUAI) Aug 2025 – present
M.Sc. in Robotics. Advisors: Xingxing Zuo, Dezhen Song Abu Dhabi, UAE

- **Research Intern** at OpenDriveLab, Shanghai Innovation Institute, Jun 2026 – Sep 2026

Tsinghua University, Yao Class (IIIS) Sep 2021 – Jun 2025
B.Eng. in Computer Science and Technology Beijing, China

- **Visiting Research Intern** at UC San Diego, Feb 2024 – Aug 2024

Publications

ForkCast: Predictive Recovery for Vision-Language-Action Policies via Counterfactual Futures

Chenru Wen*, Huajian Zeng*, Haokai Ding, Jian Kang, Anqing Duan, Dawei Zhou, Dezhen Song, Xingxing Zuo

Under review at IEEE International Conference on Robotics and Automation (ICRA), 2027.

arxiv.org/abs/2610.18650

Lay-Your-Scene: Natural Scene Layout Generation with Diffusion Transformers

Divyansh Srivastava, Xiang Zhang, He Wen, **Chenru Wen**, Zhuowen Tu

In IEEE/CVF International Conference on Computer Vision (ICCV), 2025.

arxiv.org/abs/2505.04718

Research Experience

Predictive recovery for frozen VLA policies (ForkCast) Oct 2025 – Sep 2026
RCL, MBZUAI. Advisor: Xingxing Zuo ICRA 2027, under review

- **Role.** Led all technical work: problem formulation, method, branched data, selector training, and the simulation and real-robot experiments.
- **Method.** Framed recovery as choosing among continuing the policy's chunk, moving the camera and re-querying, or replanning the arm, ranked by a candidate-conditioned latent world model and a shared progress head.
- **Simulation.** On five RoboTwin 2.0 tasks with injected occlusion and slip, frozen $\pi_{0.5}$ went from 13.3% to 30.0% pooled success (best fixed repair: 21.7%). The selector, trained only on $\pi_{0.5}$ branches, carried over to RDT and XR-0 without retraining.
- **Real robot.** On the platform below, fine-tuned $\pi_{0.5}$, collected real branches, and ran 30 trials per method in four settings (clean and restricted-visibility sponge transfer and can insertion); success rose in all four.

Localization on non-metric maps Oct 2025 – present
MBZUAI. Advisor: Dezhen Song In preparation

- Localize a walking camera on a printed visitor map, with no GPS or metric map, from local odometry and a few named landmarks: enumerate map alignments consistent with landmark bearings, then read out the nearest facility at each query.
- Evaluated on full walking sequences in Oxford; found that map-side feasibility constraints are necessary but carry little information, and that the reference itself caps attainable accuracy.

Causal analysis of appearance shifts in VLA policies Jun 2026 – Sep 2026
Research Intern, OpenDriveLab, Shanghai Innovation Institute Shanghai, China

- Turned a real-robot case, where changing only a target's appearance changed the grasp, into matched-state experiments with a fine-tuned $\pi_{0.5}$ on RoboTwin 2.0.
- Used probing and activation replacement across the visual encoder, VLM and action expert. Replacing full visual or camera-level representations recovered the donor's actions; finer interventions failed the pre-set stability bars.

Engineering

Real-robot VLA platform (dual-arm AgileX PiPER) Dec 2025 – Sep 2026
RCL, MBZUAI. Used for π_0 / $\pi_{0.5}$ deployment and all ForkCast real-robot experiments Abu Dhabi, UAE

- **Hardware.** Two PiPER arms; a RealSense head camera on an IQR PTU2 pan/tilt unit mounted on a lift that also carries both arm bases; two wrist cameras; a U-Arm leader arm for teleoperation.
- **System engineering.** Adapted openpi to the robot: a ROS 2 inference client, a Docker policy server, and real-time action chunking. Wrote the pan/tilt and lift control, a library of camera views, and the ForkCast selector service.
- **Debugging.** Recalibrated the lift's absolute height against its lower limit switch after the encoder silently re-zeroed; traced a grasp failure to a gripper command clamped at 80 mm while demonstrations opened to 100 mm.
- **Data and evaluation.** Collected demonstrations by teleoperation and replay with one-command collection and inference scripts, and ran every real-robot trial reported in ForkCast.

Honors and Awards

Lenovo Scholarship, Tsinghua University 2024
First-Class Freshman Scholarship, Tsinghua University 2021

Service

Reviewer: IEEE International Conference on Robotics and Automation (ICRA) 2027.

Skills

Programming	Python, C/C++, Bash
Learning	PyTorch; RL for dexterous hands; VLA fine-tuning (π_0, $\pi_{0.5}$); latent world models
Simulation	Isaac Lab, ManipTrans / OakInk2, RoboTwin 2.0 (SAPIEN)
Robots	AgileX PiPER dual-arm, pan/tilt and lift camera rigs, Intel RealSense